

Qwen3.8-27B Agentic Search Eval on Private Data

Han Xiao

VP of AI, Elastic

[@hxiao](#) · [in/hxiao87](#)

THIS DECK



The two serving backends

Qwen3.6-35B-A3B-MTP



hanxiao Revert README agent prompt block d501a47 · 5 days ago 37 Commits

.github/workflows	Publish prebuilt decode-optimized image to GHCR via Git...	2 months ago
docs	optimization-frontier: measured IQ4_XS speed AND qualit...	2 months ago
patches	Fix audit findings: malformed patch headers + README ra...	2 months ago
scripts	Optional pre-baked tooling image: cold start ~3.5min -> ~...	2 months ago
Dockerfile	Reconcile tok/s numbers to one signal; refresh headline fr...	2 months ago
README.md	Revert README agent prompt block	5 days ago
docker-compose.yml	Default --ctx-size 56320 everywhere (ECC off is the assu...	2 months ago

README

Qwen3.6-35B-A3B-MTP on NVIDIA L4 24GB

(Q4_K_XL) + MTP speculative decoding on a single NVIDIA L4 24GB — GCP g2-standard-8 — via the official llama.cpp Docker image. Decode-optimized to ~91-99 tok/s (min ~91 chat, max ~99 math), lossless (full GPU residency + ECC-off).

gcp gguf gpu-inference

inference-optimization

large-language-models llama-cpp

llm local-llm mixture-of-experts

moe multi-token-prediction

nvidia-l4 quantization qwen

qwen3 speculative-decoding

Readme

sparse MoE · 35B total, ~3B active · 8 of 256 experts per token · UD-Q4_K_XL ~21.3 GiB resident · MTP head baked in

91-99

tok/s decode, chat minimum to math maximum

56,320

token context

Qwen3.8-27B-UD-Q4_K_XL



hanxiao Serve on the VPC only, and add a restart path d4f3984 · 16 hours ago 8 Commits

scripts	Serve on the VPC only, and add a restart path	16 hours ago
.gitignore	Qwen3.8-27B UD-Q4_K_XL on a single NVIDIA L4	2 days ago
LICENSE	Fix repo metadata: description, topics, licence, stale com...	2 days ago
README.md	Serve on the VPC only, and add a restart path	16 hours ago

README Apache-2.0 license

Qwen3.8-27B · UD-Q4_K_XL · NVIDIA L4

Qwen3.8-27B serves at 23 tok/s decode over a 104,192-token context on a single NVIDIA L4 24 GB, using the [Unsloth Dynamic v3.0 UD-Q4_K_XL GGUF](#) and llama.cpp with MTP speculative decoding. That context is the largest the card will serve at this quantization.

Q4_K_XL GGUF) at 23 tok/s over a 104,192-token context on a single NVIDIA L4 24GB, llama.cpp + MTP speculative decoding. One-command GCP provisioning, tuning sweep, quality verification, and a binary-searched context ceiling.

benchmark gcp gguf

llama-cpp llm-inference local-llm

mtp nvidia-l4 quantization

qwen qwen3

speculative-decoding

Readme

Apache-2.0 license

Activity

0 stars

hybrid attention · 64 layers, full_attention_interval 4, so 16 hold a KV cache · native MTP head · positions to 262,144

23


tok/s decode, binary-searched ceiling

104,192

token context, the largest the card serves at this quantization

dataroom and searchbox


dataroom



A local model runs **search, read and write** in a loop until the knowledge is dumped into one fully cited .zip on disk. It stops at a measured coverage floor.

dataroom.jina.ai

searchbox



An **airgapped** testbed for search as test-time compute. One local model is locked in with a single .zip dataroom and no web access, so every answer is composed from local tools over what is in the box.

hanxiao.io/searchbox

The private corpus

GROUP	FILES	DETAIL
papers	20	six years of published work, 2023-07 to 2026-07
model-cards	29	one spec sheet per released model
blog	161	tech 104 / press 28 / insights 14 / events 8 / kb 7
api	3	OpenAPI spec, usage guide, model selection guide
faq	2	163 question-answer pairs over 14 sources
site	2	27 page namespaces, 2,297 English strings
legal	1	terms, privacy, security, DPA

Snapshot 2026-08-13, 218 files.

7,777


chunks indexed over 210 documents, median 91 units, longest 195 units at 2.1 KB

768

dimensions, jina-embeddings-v5-text-nano, L2-normalized

Six years of my own papers sit alongside the website material. Each released model carries both a paper and a release post, framed differently, and that overlap is what makes multi-hop questions possible over this corpus.

Verifier construction



A path of length k is a chain of k corpus-grounded facts, so path length sets the difficulty. The question names the starting entity and hides the bridges, which forces a solver to discover them. A question enters the set when a closed-book probe fails on it and at least one retrieval-equipped solver reaches the answer, so the set measures search over the corpus. The graphs were built over an earlier 207-file snapshot of the same material, and the 218-file snapshot is what the agentic harness indexes today.

10 graphs built over that corpus by the same local model	18,615 edges audited, every edge of every graph	40 questions sampled per graph	2 and 3 hop counts the question set covers
--	---	--	--

**A private corpus, a private verifier,
and **one low-budget GPU**
decide which model I serve.**

Han Xiao

VP of AI, Elastic

[@hxiao](#) · [in/hxiao87](#)

THE HARNESS

