

Qwen3.8-27B Agentic Search Eval on Private Data

A head-to-head replacement decision against Qwen3.6-35B-A3B, on one NVIDIA L4, over a corpus that never leaves the premises.

Han Xiao

VP of AI, Elastic

[@hxiao](#) · [in/hxiao87](#)

THIS DECK



The replacement decision

Qwen3.8 was released last week and has been widely promoted since. For agentic search on my own hardware I serve **Qwen3.6-35B-A3B** by default, and I have spent months tuning how it is served. The public reception of a new model does not tell me whether to replace it.

In agentic search over my own documents, how much better is **Qwen3.8-27B** than **Qwen3.6-35B-A3B**?

the question this evaluation answers

I evaluate two specific backends that I have already deployed and measured on the same card. A public leaderboard cannot answer the question, because the documents I search are private and the evaluation has to be built from them.

The two serving backends

Qwen3.6-35B-A3B-MTP



hanxiao

Revert README agent prompt block

d501a47 · 5 days ago

37 Commits

.github/workflows	Publish prebuilt decode-optimized image to GHCR via Git...	2 months ago
docs	optimization-frontier: measured IQ4_XS speed AND qual...	2 months ago
patches	Fix audit findings: malformed patch headers + README ra...	2 months ago
scripts	Optional pre-baked tooling image: cold start ~3.5min -> ~...	2 months ago
Dockerfile	Reconcile tok/s numbers to one signal; refresh headline fr...	2 months ago
README.md	Revert README agent prompt block	5 days ago
docker-compose.yml	Default --ctx-size 56320 everywhere (ECC off is the assu...	2 months ago

README

(Q4_K_XL) + MTP speculative decoding on a single NVIDIA L4 24GB — GCP g2-standard-8 — via the official llama.cpp Docker image. Decode-optimized to ~91-99 tok/s (min ~91 chat, max ~99 math), lossless (full GPU residency + ECC-off).

gcp

gguf

gpu-inference

inference-optimization

large-language-models

llama-cpp

llm

local-llm

mixture-of-experts

moe

multi-token-prediction

nvidia-l4

quantization

qwen

qwen3

speculative-decoding

Readme

sparse MoE · 35B total, ~3B active · 8 of 256 experts per token · UD-Q4_K_XL ~21.3 GiB resident · MTP head baked in

91-9956,320

tok/s decode, chat minimum to math maximumtoken context

Per workload: prose 92.5, code 93.0, json 92.0, chat 90.6, math 98.8, translation 91.7, summarization 93.1 tok/s. That is 45% over the 63 tok/s out-of-the-box configuration, with full GPU residency of the experts and ECC off. Raw ceiling ~73 tok/s, which MTP multiplies 1.27 to 1.37 times.

ONE L4

24 GB, 72 W, ~300 GB/s, about **\$0.24/hr spot** on a g2-standard-8. Both repositories push that one low-budget card to its measured wall, and at 91 to 99 tok/s the 35B-A3B side is comfortable inside a real agentic loop. **92 to 99 tok/s at 56K against 23 tok/s at 104K** is the throughput and window trade this evaluation resolves. Measured 2026-08-14.

Qwen3.8-27B-UD-Q4_K_XL



hanxiao

Serve on the VPC only, and add a restart path

d4f3984 · 16 hours ago

8 Commits

scripts	Serve on the VPC only, and add a restart path	16 hours ago
.gitignore	Qwen3.8-27B UD-Q4_K_XL on a single NVIDIA L4	2 days ago
LICENSE	Fix repo metadata: description, topics, licence, stale com...	2 days ago
README.md	Serve on the VPC only, and add a restart path	16 hours ago

READMEApache-2.0 license

Qwen3.8-27B · UD-Q4_K_XL · NVIDIA L4

Qwen3.8-27B serves at 23 tok/s decode over a 104,192-token context on a single NVIDIA L4 24 GB, using the [Unsloth Dynamic v3.0](#) [UD-Q4_K_XL GGUF](#) and llama.cpp with MTP speculative decoding. That context is the largest the card will serve at this quantization.

benchmark

gcp

gguf

llama-cpp

llm-inference

local-llm

mtp

nvidia-l4

quantization

qwen

qwen3

speculative-decoding

Readme

Apache-2.0 license

Activity

0 stars

hybrid attention · 64 layers, full_attention_interval 4, so 16 hold a KV cache · native MTP head · positions to 262,144


23104,192

tok/s decode, binary-searched ceilingtoken context, the largest the card serves at this quantization

At 98% of the window a 101,815-token prompt returned a correct answer at 24.88 tok/s decode and 278 tok/s prefill, with VRAM peaking at 24,082 MiB. Short-prompt prefill runs 95 to 114 tok/s. Reading 17.9 GB of weights at 300 GB/s bounds plain decode at ~16.8 tok/s, and speculation supplies the rest.

dataroom and searchbox


dataroom



A local model runs **search, read and write** in a loop until the knowledge is dumped into one fully cited .zip on disk. It stops at a measured coverage floor.

dataroom.jina.ai

searchbox



An **airgapped** testbed for search as test-time compute. One local model is locked in with a single .zip dataroom and no web access, so every answer is composed from local tools over what is in the box.

hanxiao.io/searchbox

dataroom builds the private corpus and searchbox
is the loop the two models are compared inside.

The private corpus

GROUP	FILES	DETAIL
papers	20	six years of published work, 2023-07 to 2026-07
model-cards	29	one spec sheet per released model
blog	161	tech 104 / press 28 / insights 14 / events 8 / kb 7
api	3	OpenAPI spec, usage guide, model selection guide
faq	2	163 question-answer pairs over 14 sources
site	2	27 page namespaces, 2,297 English strings
legal	1	terms, privacy, security, DPA

Snapshot 2026-08-13, 218 files.

7,777


chunks indexed over 210 documents, median 91 units, longest 195 units at 2.1 KB

768

dimensions, jina-embeddings-v5-text-nano, L2-normalized

Six years of my own papers sit alongside the website material. Each released model carries both a paper and a release post, framed differently, and that overlap is what makes multi-hop questions possible over this corpus.

Verifier construction



A path of length k is a chain of k corpus-grounded facts, so path length sets the difficulty. The question names the starting entity and hides the bridges, which forces a solver to discover them. A question enters the set when a closed-book probe fails on it and at least one retrieval-equipped solver reaches the answer, so the set measures search over the corpus. The graphs were built over an earlier 207-file snapshot of the same material, and the 218-file snapshot is what the agentic harness indexes today.

10 graphs built over that corpus by the same local model	18,615 edges audited, every edge of every graph	40 questions sampled per graph	2 and 3 hop counts the question set covers
--	---	--	--

Evaluation design

		CLOSED-BOOK	SINGLE-SHOT BM25	AGENTIC BM25, 4 ROUNDS
Qwen3.8-27B	104,192 ctx	control that removes items answerable from memory	top-5 passages, one shot	the condition under test
Qwen3.6-35B-A3B	56,320 ctx	control, same corpus	top-5 passages, one shot	the incumbent

HELD FIXED

The corpus, the question set, the retrieval index, the grading procedure and the four-round cap. The agent runs in the airgapped loop with one retrieval tool over embeddings.npz and chunks.jsonl, plus stock file access. Only the served model changes across rows.

SCORING

The answer held at the end of every turn is captured and graded by the same local judge, so accuracy is reported against turn index and against cumulative fresh-prefill input tokens. Turns are what a person reads and fresh-prefill tokens are what the GPU pays.

**A private corpus, a private verifier,
and **one low-budget GPU**
decide which model I serve.**

Han Xiao

VP of AI, Elastic

[@hxiao](#) · [in/hxiao87](#)

THE HARNESS

