

— ZERO TRAINING · NO TASK HEAD · OPEN VOCABULARY

Test-time compute for image tagging with a frozen jina-embeddings-v5-omni

Han Xiao

VP of AI, Elastic

[@hxiao](#) · [in/hxiao87](#)

GITHUB REPO



the code

The previous talk: search as test-time compute over a frozen encoder.

The claim from the autoresearch talk: you do not need a bigger model to get better retrieval. You **recombine the geometry** of one frozen encoder at inference: split the document, re-score, fuse channels, feed back. That extra work at test time is the compute.

THE MODEL

frozen. One embedding API. No retraining, no new parameters.

THE COMPUTE

multi-pass embedding algebra assembled at test time, climbing a cost-versus-quality curve.

Same principle, new question: beyond buying more **relevance** on the trained task, can test-time compute produce an entirely **new capability**?

Test-time compute for an embedding model: three families.

Language models scale test-time compute in tokens: longer chains, more samples. A frozen encoder has no tokens to spend. Its inference-time budget takes exactly three forms:

A MORE COMPUTATION PER PASS

Extract more from a single forward pass: score every patch, fuse channels, subtract priors. Cost: a few matrix multiplies.

LLM analogue: a richer readout of one sample.

B NEW PASSES, NEW VIEWS

Run the frozen encoder **again** on inputs it has not seen: crops of the image, splits of the document. Cost: full forward passes.

LLM analogue: best-of-N sampling.

C RE-PROCESS THE FEATURES

Transform the vectors already computed: whitening, graph propagation, optimal transport, re-normalization. Cost: nearly zero.

Where most of the 2024–2026 training-free literature operates.

This talk is a controlled experiment across all three families.
Two of them scale accuracy with compute. One does not.

Apply the same principle to a task the model was never trained for.

jina-embeddings-v5-omni-nano is a ~1B retrieval model. It embeds text and images into one space for search. It has **no tagging head, no classifier, no tag list**. It was never trained to answer "what is in this image".

THE STANDARD APPROACH

train a tagging head. Curate a label set. Fine-tune. A second model, more data, more parameters.

THIS TALK

keep the model frozen. Construct the entire tagger from test-time compute over its existing geometry.

If a capability the model was never trained for **emerges** purely from inference-time computation, then test-time compute yields new capability, not merely higher quality.

— THE TASK

Open-vocabulary, multi-label image tagging.

Give it an image, get back the words for what is in it. Not one label: **every** object present. Not a fixed 80-class head: an **open vocabulary**, any word the model knows.

- | OPEN-VOCABULARY no curated tag list; the label space is the model's own.
- | MULTI-LABEL a photo has many objects; recall them all.
- | SINGLE FROZEN MODEL no WordNet, no POS tagger, no second network.



cat.jpg (86ms) → [kitty](#), [cosy](#), [plush](#), [crib](#), [paw](#), [sleeps](#)

Every signal must come from the one frozen model.

ZERO TRAINING

NO SECOND MODEL


NO WORDNET / DICTIONARY

NO REGEX / POS TAGGER

This is the same discipline as the retrieval talk, applied to a new task. The label space, the visual features, and the word filter must **all** come from the frozen encoder. The moment a curated tag list or a POS tagger is added, it is no longer the model doing the work.

The constraint is the experimental control: it ensures the entire capability is attributable to **test-time compute over existing geometry**.

One frozen encoder, two input paths, one 768-d space.



Everything lands in one space: **any patch of any image, scored against any word the model knows** — $\langle p, e_\ell \rangle$ is a meaningful number.


— THE PIPELINE

The complete tagger: test-time compute around a frozen forward pass.

OFFLINE · COMPUTED ONCE, CACHED



PER IMAGE · 75 MS FAST · 1016 MS --HQ



Every box is the frozen model or plain arithmetic. **A** is more algebra on one pass; **B** is new passes on new pixels. No box is family C.

The label set is the tokenizer's own vocabulary.

No curated tag list. Take all **128,260** tokens and encode each one as `text` into a label matrix

$$E = [e_\ell]_{\ell \in V}, \quad e_\ell = \text{encode_text}(\ell) \in \mathbb{R}^{768}, \quad |V| = 128,260$$

One matmul now scores any image against the entire vocabulary.

```
wrong: image_vec · embed_tokens.weight  
       → Tutor, avatar, PyTuple (garbage)  
right: image_vec · encode_text(token)  
       → kitty, cosy, plush, paw (aligned)
```

Why `encode_text`, not the embedding table? The model has `tie_word_embeddings = false`: the input-embedding table lives in a different space than the pooled output. You must encode each token *as a text string* to put labels in the image's space.

The original intuition — use the vocabulary directly — was correct, but only through the aligned output space.

One pooled vector under-detects; score every patch instead.

GLOBAL POOLED




$$s_\ell = \langle g, e_\ell \rangle$$

one vector dominated by the most salient object

MAP 0.264

PATCH MAX + GLOBAL




$$s_\ell = 0.7 \max_{p \in P} \langle p, e_\ell \rangle + 0.3 \langle g, e_\ell \rangle$$

a small object fires on its own patch and survives

MAP 0.635

THE GAIN



+0.37 mAP

"[CLS] is not enough", confirmed on this encoder

SAME PIXELS, MORE COMPUTATION


A Pure family-A compute: the patch features already exist in the forward pass. **One additional max over rows yields +0.37 mAP.**

Generic words sit close to every image: subtract the per-label prior.

Raw cosine has a base-rate bias: generic words ("bed", "cat") score high on almost any image because they sit near the image modality. Estimate a per-label prior once from neutral background images, then subtract it:

$$\tilde{s}_\ell = s_\ell - \mu_\ell, \quad \mu_\ell = \mathbb{E}_{x \sim \text{background}} [s_\ell(x)]$$


Pure marginal de-biasing: no calibration set, no labels, one offline pass. A word that scores high on *everything* gets centered to zero; only image-specific spikes survive. It is the lightest possible correction, and it is enough because the space is already well aligned.



The tokenizer already knows what a word is.

Word-start gate. Byte-BPE marks whole-word starts with a space glyph. `Ġcat` is a word; `GetComponent` is a fragment. That one built-in signal filters 128k tokens down to **25,465** clean words. No external dictionary.

Embedding-NMS. Synonyms and multilingual variants (/ `Cat` / `кот` / `cat`) collapse into one, using cosine in the model's own space. Walk labels by score; keep ℓ only if $\max_{k \in \text{kept}} \langle e_\ell, e_k \rangle < 0.6$. Again: no lookup table, just the geometry.




CWR multi-crop: the one lever that raises the accuracy ceiling.

Re-encode the image as a **3×3 + 2×2 + center** grid of 14 crops through the same frozen model, and take the per-label **max** across crops. A small object becomes large in whichever crop contains it, so its signal jumps from weak to strong.



$$S_\ell = \tilde{s}_\ell + 1.3 \left(\max_{c=1..14} s_{c,\ell} - \mu_\ell \right), \quad s_{c,\ell} = \max \left(\max_{p \in P_c} \langle p, e_\ell \rangle, \langle g_c, e_\ell \rangle \right)$$

B Test-time augmentation done right: full-coverage grid plus per-label **max**. Blind center-crop averaging **hurt**; the outlier crop is not noise, it is the evidence.



fast: wolf, roar, muzzle → --hq: fur, bear, muzzle
CWR recovers the small bear the single pass misses

The entire tagger is one scoring function.

$$S(\ell) = \underbrace{0.7 \left(\max_{p \in P} \langle p, e_\ell \rangle - \mu_\ell \right)}_{\text{patch evidence}} + \underbrace{0.3 \left(\langle g, e_\ell \rangle - \mu_\ell^g \right)}_{\text{global context}} + \underbrace{1.3 \left(\max_{c=1..14} s_{c,\ell} - \mu_\ell \right)}_{\text{multi-crop re-encode}}$$

A Patch evidence — algebra over rows the forward pass already produced. Free. +0.37 mAP.

A Global context — the pooled vector, de-biased by its own background prior. Free.

B Multi-crop — 14 fresh forward passes on new pixels. 14× the cost. +0.075 mAP.

then $\ell \in \hat{G}$ -gated words $\rightarrow$ embedding-NMS ($\tau = 0.6$) $\rightarrow$ top- k — and this composition of arithmetic is the tagger.


No head, no logits, no learned threshold. Every symbol is either the frozen model or a max.

RESULTS

COCO-150, 80 categories: a frozen retrieval model that tags.


METHOD	P@1	P@3	R@5	MAP
global pooled	0.433	0.289	0.449	0.264
patch max	0.753	0.427	0.631	0.635
softpool	0.773	0.449	0.649	0.608
+ CWR multi-crop	0.813	0.476	0.680	0.710

150 COCO-val images, real multi-label ground truth, avg 2.93 labels/image.



mAP: global → patch → +CWR

Accuracy as a function of test-time compute.



The first gain is **nearly free**: the same pixels, more computation per pass. The second costs 14× the latency for +0.075. The pink cluster spent its compute re-processing *features* — and gained nothing.

Open vocabulary, out-of-distribution photos.



AI conference hall

onstage

attendees

ceilings

seats



Porsche 911 interior · --hq

carro

leather

steering

clock

seat



bear on grass · --hq

fur

bear

muzzle

ivory

Labels come directly from the tokenizer, so multilingual variants ("carro") and occasional fragments appear. That is the direct cost of a **zero-annotation** open vocabulary.

Does accuracy continue to scale with test-time compute?

The tagger works. The scientific question is the same one from the retrieval talk: **does accuracy scale with the compute you pour in at test time?** We systematically implemented the obvious upgrades from the 2024-2026 training-free literature and measured each on COCO-150.

C RE-PROCESS THE FEATURES

change the layer, whiten, softmax over classes, propagate on a graph, optimal transport, Bayesian priors, Dirichlet, robust aggregation.


B FEED THE MODEL NEW INFO

CWR: run the frozen encoder again on crops it has not seen at that scale.

One of these two families works. The other does nothing, or collapses.

— EVERYTHING WE TRIED

Nine training-free methods from the literature; one beats the baseline.



Re-processing well-calibrated features yields nothing; new information does.

Every method that **re-processes** the existing features - new layer, whitening, cross-class softmax, label propagation, optimal transport, Dirichlet, robust trimming - is a no-op or a collapse. They were designed for weakly calibrated single-label CLIP. This encoder's space is already good and already multi-label, so there is nothing left to recover.

RE-PROCESSING FEATURES

whitening 0.06 · ZLaP 0.14 · EM-Dirichlet 0.17 · softmax 0.636 · OTTER 0.699. All $\leq$ baseline.

FEEDING NEW INFORMATION


CWR crops $\rightarrow$ mAP 0.710, P@1 0.813. The only lever that moves the ceiling.

The ceiling is the 1B model itself. Real test-time compute here means **giving the model more to look at**, not re-arranging what it already saw.

From study to production: on-device tagging in Omni.

The Python study became a feature of **Omni**, a native on-device semantic-search app (Swift + MLX, same frozen model). The port is where the design pays off:

- SAME FORWARD PASS** the patch rows already exist for the file's *embedding*; tagging adds one matmul against a memory-mapped bf16 label matrix.
- ~0.6 MS / IMAGE** about 4% overhead on the embed step. Tagging is effectively free at index time.
- TAGS = SNIPPETS** tags land in the search snippet, so photos and videos become findable by *keyword*, not just by vector.
- SELF-CALIBRATING** the background prior μ is estimated on-device from the first 64 images it sees, then frozen.



tag chips as they appear in results; the third row is a video segment

One label matrix: images, high-quality re-passes, video, and scanned documents.



IMAGE

Tagged at **index time**, inside the embedding forward. Patch-max + global, prior-centered, gated, NMS-deduped.

+1 matmul · ~4% overhead



HQ IMAGE

A background **re-tag pass** adds 5-crop CWR: 2x2 grid with 15% overlap, plus the center. Per-label max, fused at weight 1.0.

P@1 0.773 → 0.847 · 14 crops =
+0.01 mAP for 3× cost, not shipped



VIDEO

Up to 32 frames sampled per **240-second segment** run through the tower as one sequence. The patch-max pools over **space and time**: a bear in one frame of one segment still fires.

one tag set per segment



SCANNED PDF

Non-extractable pages render to images and take the same path: **per-page tags** like "invoice, table, signature" become the page's snippet.

every page, keyword-findable

One label matrix, one scoring rule. The media shape only changes **what counts as a patch**: crops for detail, frames for time, pages for documents.

A capability the model was never trained for, realized at inference.

Two talks, one thesis. A frozen encoder holds more capability than its training objective exposes, and you unlock it with **test-time compute**, not more parameters.

RETRIEVAL TALK

recombine the geometry for more **relevance** on the task it was trained for.

THIS TALK

recombine the geometry for an **emergent new task** it was never trained for: tagging.

Test-time compute scales only when it **supplies the model with new information**. Re-processing an already-good representation yields nothing.

Versus the 2024-2026 training-free tagging literature.

LINE OF WORK	THEIR SETUP	OUR DELTA
RAM / RAM++	trained tagging model, curated 6.4k tag list	zero training, label space = tokenizer vocab
TagCLIP	CLIP two-tower, penultimate layer, DMAR+CWR	single omni model, last layer is the output space
PIAA	GDA whitening to close the modality gap	image/text already one space; whitening collapses it
2025 calibration (OTTER, BCA)	calibration set or target prior to de-bias	one offline background prior, zero calibration

The novelty is the **combination**: tokenizer vocab as the open label space, one frozen omni model for every signal, and marginal de-biasing, with all the fancy test-time machinery empirically ruled out.

**A frozen model already knows more
than its objective admits.**
Test-time compute is how you ask.

Han Xiao · VP of AI, Elastic

[@hxiao](#) · [in/hxiao87](#)

github.com/hanxiao/jina-v5-omni-nano-test-time-image-tagging

GITHUB REPO



the code